

Linear Regression Algorithms

Kashyap Sureshchandra Patel

September 7, 2025

Algorithm 1: Linear Regression with Gradient Descent and Regularization

Input: Learning rate η , number of epochs T , batch size b , regularization flag r , regularization method $m \in \{\text{L1, L2, ElasticNet}\}$, regularization strength λ , mixing parameter α

Output: Optimal parameter vector θ

Initialize: $\theta \leftarrow \mathbf{0}$

Add bias column of ones to X : $X_b = [\mathbf{1}, X]$

$n \leftarrow$ number of samples, $d \leftarrow$ number of features (including bias)

for $epoch = 1$ to T **do**

 Shuffle dataset (X_b, y) randomly

for each mini-batch (X_{batch}, y_{batch}) of size b **do**

$y_{pred} \leftarrow X_{batch} \cdot \theta$

 Gradient: $\nabla J(\theta) \leftarrow \frac{2}{b} X_{batch}^\top (y_{pred} - y_{batch})$

if $r = \text{True}$ **then**

$\theta_{reg} \leftarrow \theta$, set $\theta_{reg,0} \leftarrow 0$

if $m = \text{L1}$ **then**

$\nabla J(\theta) \leftarrow \nabla J(\theta) + \lambda \cdot \text{sign}(\theta_{reg})$

else if $m = \text{L2}$ **then**

$\nabla J(\theta) \leftarrow \nabla J(\theta) + 2\lambda \cdot \theta_{reg}$

else if $m = \text{ElasticNet}$ **then**

$\nabla J(\theta) \leftarrow \nabla J(\theta) + \lambda \left(\alpha \cdot \text{sign}(\theta_{reg}) + 2(1 - \alpha) \cdot \theta_{reg} \right)$

 Update step: $\theta \leftarrow \theta - \eta \cdot \nabla J(\theta)$

Algorithm 2: Linear Regression with Normal Equation (OLS and Ridge)

Input: Training data (X, y) , regularization method $m \in \{\text{None}, \text{L2}\}$, regularization strength λ

Output: Optimal parameter vector θ

Preprocessing:

Add bias column of ones to X : $X_b = [\mathbf{1}, X]$

$n \leftarrow$ number of samples, $d \leftarrow$ number of features (including bias)

if $m = \text{None}$ (*OLS*) **then**

$\theta \leftarrow (X_b^\top X_b)^\dagger X_b^\top y$

else if $m = \text{L2}$ (*Ridge*) **then**

$I_d \leftarrow d \times d$ identity matrix

 Set $I_{d,0,0} \leftarrow 0$ (do not regularize bias)

$\theta \leftarrow (X_b^\top X_b + n\lambda I_d)^\dagger X_b^\top y$

return θ
